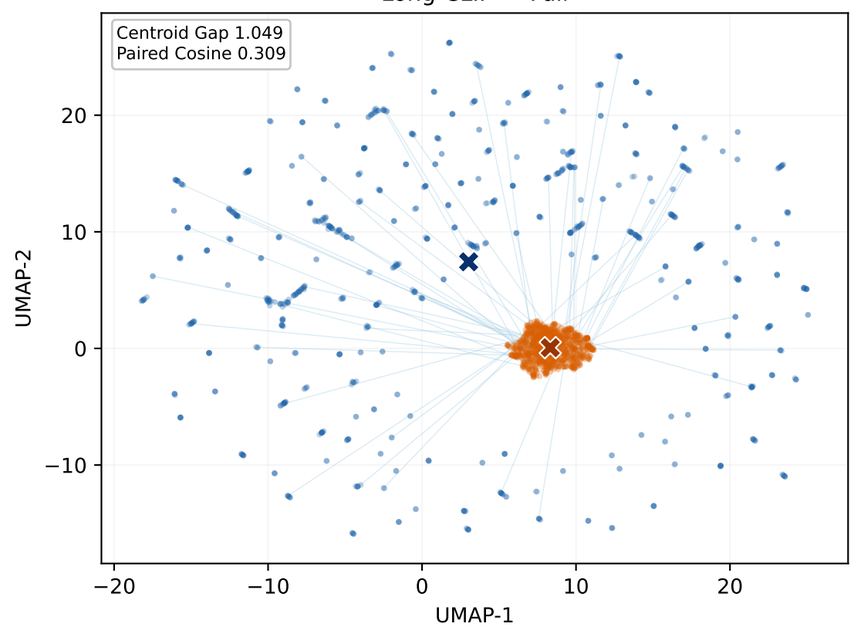
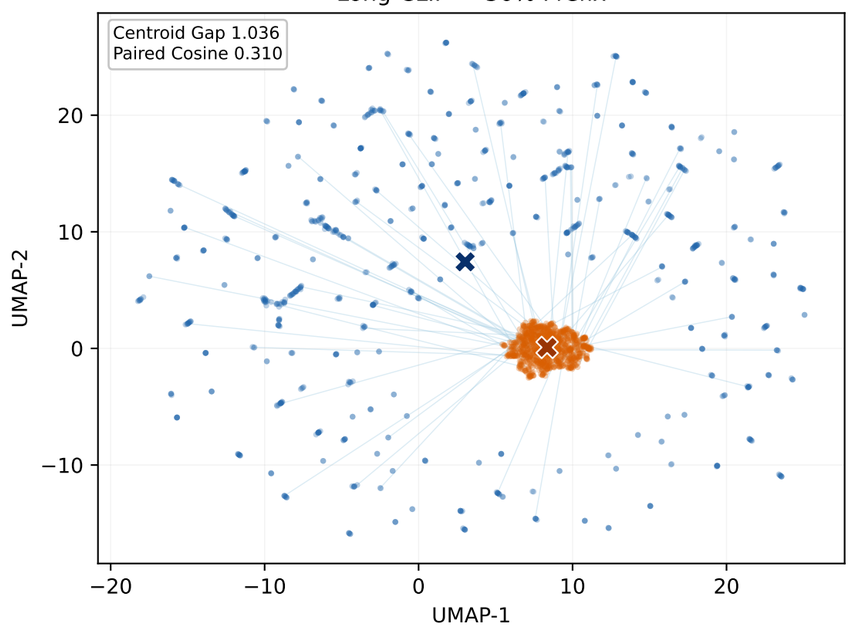
DOCCI modality gap — model-specific shared UMAP fits

Image · Text · Image centroid · Text centroid · 50 paired links

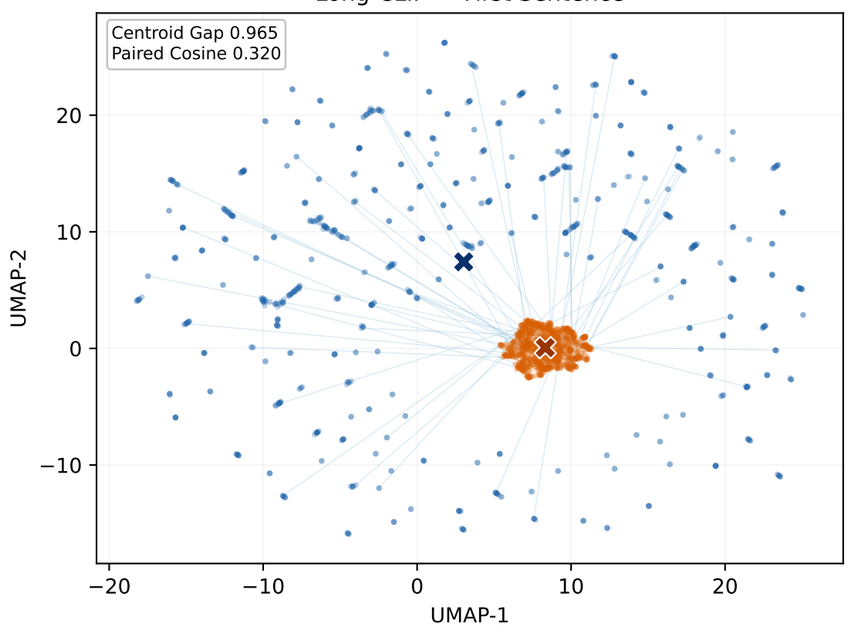
Long-CLIP — Full



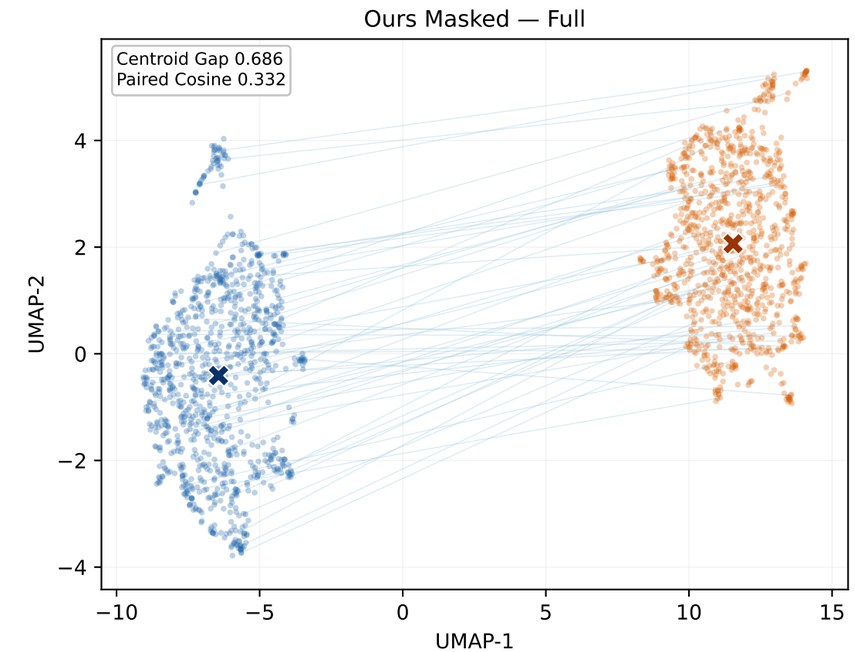
Long-CLIP — 50% Prefix



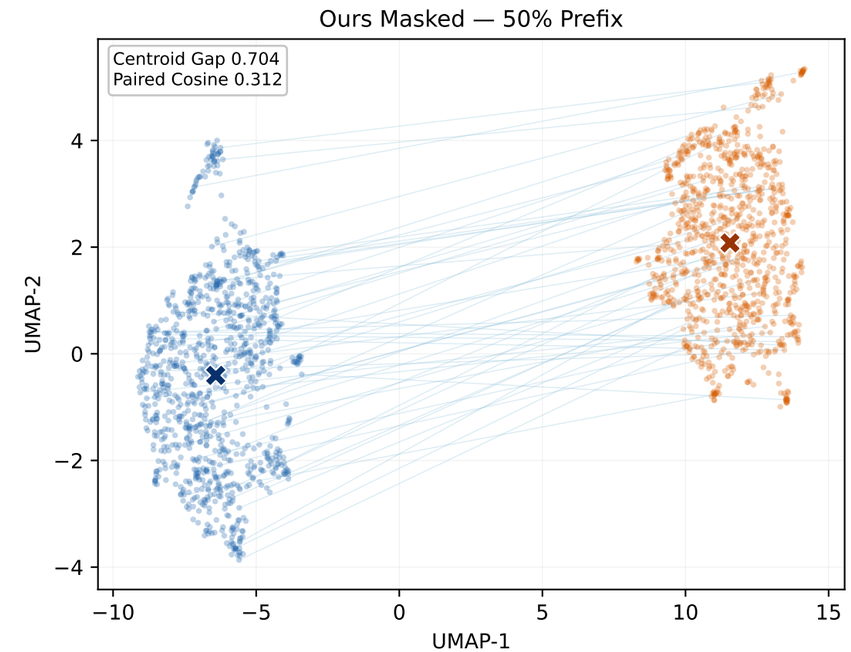
Long-CLIP — First Sentence



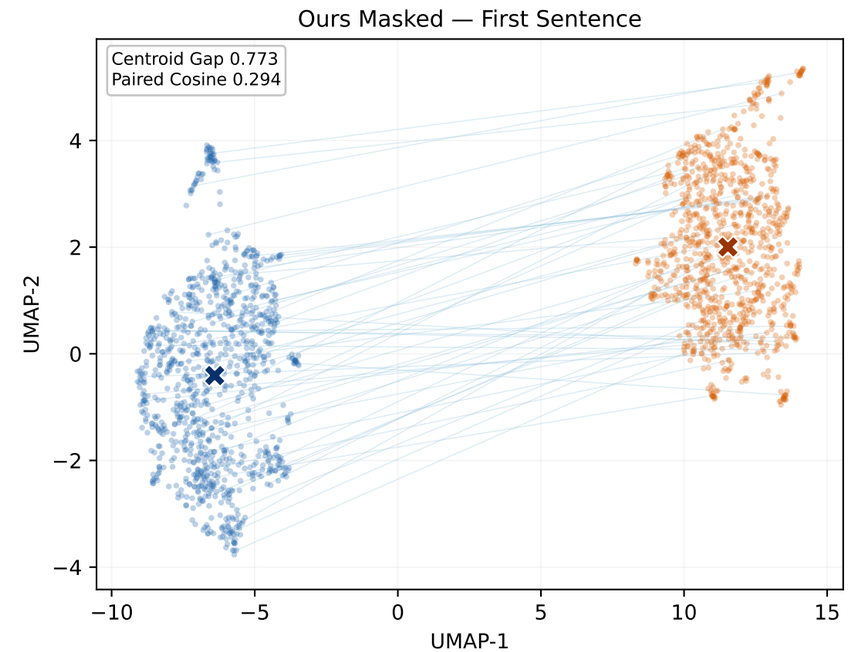
Ours Masked — Full



Ours Masked — 50% Prefix



Ours Masked — First Sentence



Each model row uses one UMAP fit over all three caption settings; rows are not on a shared UMAP.